

# Introduction to Deep Learning

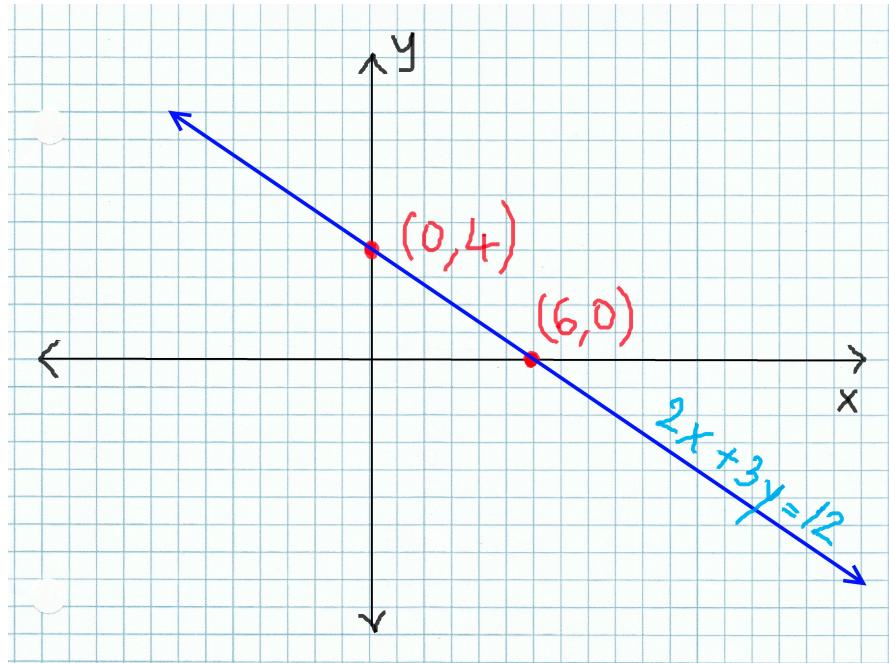
## 4. Linear Regression, Basic Optimization

STAT 157, Spring 2019, UC Berkeley

Alex Smola and Mu Li

[courses.d2l.ai/berkeley-stat-157](https://courses.d2l.ai/berkeley-stat-157)

# Linear Methods



# House Buying 101

- Pick a house, take a tour, and read facts
- Estimate its price, bid it

Listing price  
from agent

\$5,498,000 | 7 | 5 | 4,865 Sq. Ft.  
Price Beds Baths \$1130 / Sq. Ft.  
Redfin Estimate: \$5,390,037 On Redfin: 15 days

Predicted  
sale price



## Virtual Tour

- [Branded Virtual Tour](#)
- [Virtual Tour \(External Link\)](#)

## Parking Information

- Garage (Minimum): 2
- Garage (Maximum): 2
- Parking Description: Attached Garage, On Street
- Garage Spaces: 2

## Multi-Unit Information

- # of Stories: 2

## School Information

- Elementary School: El Carmelo Elementary School
- Elementary School District: Palo Alto Unified School District
- Middle School: Jane Lathrop Stanford Middle School
- High School: Palo Alto High School
- High School District: Palo Alto Unified School District

## Interior Features

### Bedroom Information

- # of Bedrooms (Minimum): 7
- # of Bedrooms (Maximum): 7

- Kitchen Description: Countertop (Granite), Dishwasher, Garbage Disposal, Hood Over Range, Island with Sink, Microwave, Oven Range

# House Price Predicting

- Very important, that's real money...



\$100K+ gap

Redfin over-estimated the price, and we believed it

# A Simplified Model

- Assumption 1: the key factors impacting the prices are
  - #Beds, #Baths, Living Sq. Ft.
  - Denote by  $x_1, x_2, x_3$
- Assumption 2: the sale price is a weighted sum over the key factors



$$y = w_1x_1 + w_2x_2 + w_3x_3 + b$$

- Weights and bias are determined later

# Linear Model

- Given  $n$ -dimensional inputs  $\mathbf{x} = [x_1, x_2, \dots, x_n]^T$
- Linear model has a  $n$ -dimensional weight and a bias

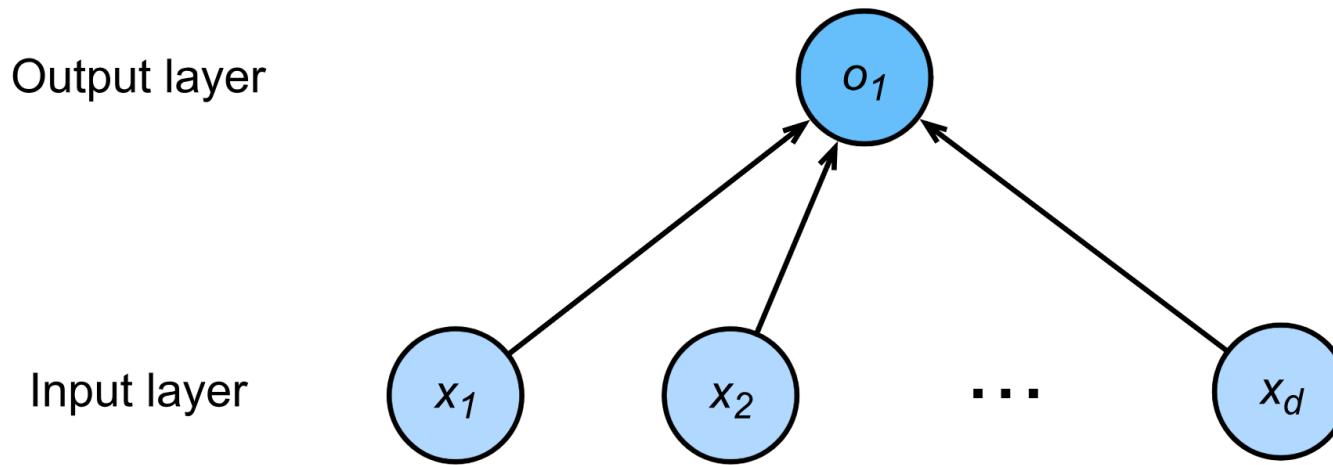
$$\mathbf{w} = [w_1, w_2, \dots, w_n]^T, \quad b$$

- The output is a weighted sum of the inputs

$$y = w_1x_1 + w_2x_2 + \dots + w_nx_n + b$$

- Vectorized version  $y = \langle \mathbf{w}, \mathbf{x} \rangle + b$


# Linear Model as a Single-layer Neural Network



- We can stack multiple layers to get deep neural networks

# Neural Networks Derive from Neuroscience

## The real neuron



# Measure Estimation Quality

- Compare the true value vs the estimated value
  - Real sale price vs estimated house price
- Let  $y$  the true value, and  $\hat{y}$  the estimated value, we can compare the loss

$$\ell(y, \hat{y}) = (y - \hat{y})^2$$

- It is called squared loss

# Training Data

- Collect multiple data points to fit parameters
  - Houses sold in the last 6 months
- It is called the training data
- The more the better
- Assume  $n$  examples

$$\mathbf{X} = [\mathbf{x}_0, \mathbf{x}_1, \dots, \mathbf{x}_n]^T \quad \mathbf{y} = [y_0, y_1, \dots, y_n]^T$$

# Learn Parameters

- Training loss

$$\ell(\mathbf{X}, \mathbf{y}, \mathbf{w}, b) = \frac{1}{n} \sum_{i=1}^n (y_i - \langle \mathbf{x}_i, \mathbf{w} \rangle - b)^2 = \frac{1}{n} \| \mathbf{y} - \mathbf{X}\mathbf{w} - b \| ^2$$

- Minimize loss to learn parameters

$$\mathbf{w}^*, \mathbf{b}^* = \arg \min_{\mathbf{w}, b} \ell(\mathbf{X}, \mathbf{y}, \mathbf{w}, b)$$

# Closed-form Solution

- Add bias into weights by  $\mathbf{X} \leftarrow [\mathbf{X}, \mathbf{1}] \quad \mathbf{w} \leftarrow \begin{bmatrix} \mathbf{w} \\ b \end{bmatrix}$

$$\ell(\mathbf{X}, \mathbf{y}, \mathbf{w}) = \frac{1}{n} \| \mathbf{y} - \mathbf{X}\mathbf{w} \|^2 \quad \frac{\partial}{\partial \mathbf{w}} \ell(\mathbf{X}, \mathbf{y}, \mathbf{w}) = \frac{2}{n} (\mathbf{y} - \mathbf{X}\mathbf{w})^T \mathbf{X}$$

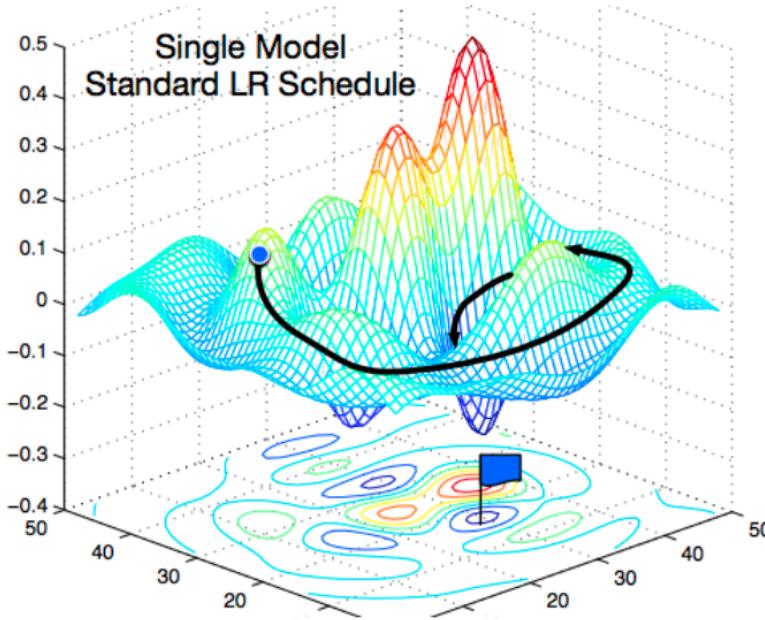
- Loss is convex, so the optimal solutions satisfies

$$\frac{\partial}{\partial \mathbf{w}} \ell(\mathbf{X}, \mathbf{y}, \mathbf{w}) = 0$$

$$\Leftrightarrow \frac{2}{n} (\mathbf{y} - \mathbf{X}\mathbf{w})^T \mathbf{X} = 0$$

$$\Leftrightarrow \mathbf{w}^* = (\mathbf{X}^T \mathbf{X})^{-1} \mathbf{X} \mathbf{y}$$


# Basic Optimization



# Gradient Descent

- Choose a starting point  $\mathbf{w}_0$
- Repeat to update the weight  $t=1, 2, 3$


$$\mathbf{w}_t = \mathbf{w}_{t-1} - \eta \frac{\partial \ell}{\partial \mathbf{w}_{t-1}}$$




- Gradient: a direction that increases the value
- Learning rate: a hyper-parameter specifies the step length

# Choose a Learning Rate

Not too small



Not too big



# Mini-batch Stochastic Gradient Descent (SGD)

- Computing the gradient over the whole training data is too expensive
  - Takes minutes to hours for DNN models
- Randomly sample  $b$  examples  $i_1, i_2, \dots, i_b$  to approximate the loss

$$\frac{1}{b} \sum_{i \in I_b} \ell(\mathbf{x}_i, y_i, \mathbf{w})$$

- $b$  is the batch size, another important hyper-parameters

# Choose a Batch Size

Not too small

Workload is too small, hard  
to fully utilize computation  
resources

Not too big

Memory issue  
Waste computation, e.g.  
when all  $x_i$  are identical

# Summarize

- Problem: estimate a real value
- Model:  $y = \langle \mathbf{w}, \mathbf{x} \rangle + b$
- Loss: square loss  $\ell(y, \hat{y}) = (y - \hat{y})^2$
- Learn by mini-batch SGD
  - Choose a starting point
  - Repeat
    - Compute gradient
    - Update parameters